

The cognitive turn in experimental pragmatics: From common ground to heuristic predictions of the dynamics of resource-efficient minds

Michael Franke

Abstract

concrete.

1 Experimental pragmatics and computational modeling

The linguistic subfield of pragmatics studies how language is used for efficient communication and how communicated meaning is modulated, if not created, in a particular context by systematic, context-general patterns of use. Its origins lie in the philosophy of language, in particular the analytical tradition of philosophy, in work on topics like *speech-acts* (Austin, 1962; Searle, 1969), *conversational implicature* (Grice, 1975), or *presupposition* (Stalnaker, 1973; Karttunen, 1973). Much early work on these topics remained invested in a logical-formal approach (Gazdar, 1979; Kadmon, 2001), but the early 2000’s brought a deep and empowering incision: a widely adopted and ultimately very successful *experimental turn* (Noveck and Sperber, 2004; Noveck, 2018), in parallel to similar empirical re-orientations in other hitherto very analytically minded subfields of linguistics like formal syntax (Sprouse, 2007, 2023) and formal semantics (Meibauer and Steinbach, 2011; Cummins and Katsos, 2019).

The experimental turn in pragmatics was accompanied by formal and computational work from the start, with two main strands of approaches emerging at opposite ends of a spectrum. On one end of the spectrum, *grammaticalism* focuses on language as an abstract, computationally encapsulated system [MF: refs]. Grammaticalism strives to give formally precise explanations for how pragmatic inferences (focusing almost exclusively on interpretation) could arise in a language-specific computational module [MF: refs]. On the other end of the spectrum, *probabilistic pragmatics* seeks explanations of pragmatic phenomena from the point of view that linguistic communication as a special form of social interaction, also taking aspects of domain-general cognition of language users into account [MF: refs]. Figure 1 gives a coarse-grained overview of the development of, in particular, probabilistic pragmatics as the culmination point of research that combined aspects of individual cognition, such as bounded rationality [MF: refs], and functional explanations based on optimal information transfer in communication [MF: refs].

While experimental data can mitigate between alternative explanations proposed in the grammaticalist tradition [MF: refs], I argue here that the more socio-cognitive approach adopted by probabilistic pragmatics is, if done right, much better suited to incorporate *both* linguistic theory *and* aspects of general cognition (such as perception, adaptation and decision-making), as well as




Figure 1: Genealogy of probabilistic pragmatics. [MF: describe what’s pictured here (if you keep it; maybe make this figure and the text part of an InfoBox?)]

the concrete affordances of a conversational situation or an experimental task, into a unified computational model of the *data-generating process*, i.e., the (cognitive) processes by which individual data observations, such as concrete data points recorded from an experiment, come about. Seen in this light, computational modeling in pragmatics can serve a bridging function between language, linguistic cognition and domain-general aspects of information-processing and decision-making. The central argument of this position paper is, consequently, that a possible and highly productive future of experimental pragmatics in general consists in a dedicated *cognitive turn*, which seeks to explain language use as the intricate interplay between language-specific and domain-general cognition. Computational cognitive-pragmatic models plays a pivotal role in this cognitive turn, as they require detailed commitment about the role of linguistic theory (latent, established concepts) in the data-generating process, and the way that general cognitive factors shape linguistic behavior. As I will argue here, recent work in modelling-oriented experimental pragmatics has started to “go cognitive” in this sense, but that a more dedicated embrace of the cognitive turn is needed. Looking forward, I sketch possible future directions.

[MF: spell out paper overview] Section [MF: XYZ] introduces probabilistic pragmatics models. Section [MF: XYZ] elaborates on the idea of cognitive models acting as theory-driven statistical models for the data-generating process. It covers some main lessons learned from probabilistic pragmatic modeling of experimental data, and points to key open questions for the future. Section [MF: XYZ] then focuses on a high-level case study, namely probabilistic reasoning about common ground, as an illustrative example ...

2 Probabilistic pragmatics

- Notable early probabilistic approaches to pragmatics include pioneering work by Merin (1994, 1999), Zeevat (2014), and Russell (2012)
- particularly prominent style of probabilistic modeling in pragmatics is the Rational Speech Act framework (Frank and Goodman, 2012; Scontras, Tessler, and Franke, 2021; Degen, 2023), the basic set-up of which has been anticipated in decision-theoretic work by Anton Benz (Benz, 2006; Benz and van Rooij, 2007)
- what unites probabilistic pragmatics approaches is, arguably, the idea that pragmatic interpretation is a form of inverse reasoning, or abductive reasoning Hobbs, Stickel, and Martin (1993), Bunt and Black (2000), and Franke and Jäger (2016): since it is natural to assume that the speaker’s linguistic behavior (their choice of words, that they spoke in the first place, their intonation or speech-accompanying gesture, etc.) utterance is at least partly defined by a latent, i.e., unobservable, mental state of the speaker, the main task of pragmatic interpretation is to reason backwards, as it were, about how relatively likely or unlikely different conceivable causes for the speaker’s behavior are
- This type of reasoning is modelled by Bayesian probabilistic reasoning: based on a probabilistic speaker policy $\pi(a | t)$ that determines how likely that speaker is to execute an action $a \in A$ if they are in state $t \in T$, and a prior distribution about how likely different states t are, the normatively correct way of updating these prior beliefs is given by Bayes rule as follows:

$$P(t | a) = \frac{\pi(a | t) P(t)}{\sum_{t' \in T} \pi(a | t') P(t')} \quad (1)$$

- to fill this basic (and arguably natural (at least as a starting point)) assumption of interpretation with life, at least three things have to be specified:
 - (i) the relevant set of states T ,
 - (ii) the relevant set of actions A , and
 - (iii) the speaker’s behavioral policy π .
- For a standard case of referential communication, for instance, the states T could be potential referents in the shared perceptual context between speaker and listener, and the actions A could be a finite set of utterances or descriptions, e.g., identified at the level of surface form, such as “the blue triangle.” The speaker’s policy π could then be defined based on standard assumptions about how speakers are expected to behave in situations of the relevant kind, e.g., as envisaged by Grice’s Maxims of Conversation (Grice, 1975). In particular, for a simple case of referential communication, we might define a probabilistic policy in terms of a utility function that score each action (utterance) a in a given state t based on a ’s truth-value in t (quality), the degree to which a conveys information about t in distinction to any other $t' \in T$, and the cost $C(a)$ of performing a :

$$U(a, t) = \text{Informativity}(a) + \text{Truth}(a, t) + \text{Cost}(a)$$

The speaker’s policy is then defined as a, possibly approximate, optimization of this utility function, e.g., as implemented by the softmax operation (Franke and Degen, 2023):

$$\pi(a \mid t) = \frac{\exp \alpha U(a, t)}{\sum_{a' \in A} \exp \alpha U(a', t)}$$

- The standard formulation of the RSA-speaker policy, as given by Frank and Goodman (2012) and much subsequent work, adheres to this general constraint-optimization picture, but is usually given in terms of a speaker optimizing their linguistic expressions for a literal listener. But as the literal listener is usually constructed as a context-blind, unstrategic dummy, we should more generally consider a probabilistic pragmatic model to consist of:
 - a speaker policy that seeks to approximate optimal behavior under some (numeric) utility defining “efficient / successful communication;” and
 - a Bayesian listener who inverts the speaker policy to infer likely explanations of the observed linguistic behavior.
- [MF: maybe include:] higher-order iteration; multiple unknowns; more factors in the speaker utility function

3 Theory-driven probabilistic modeling

- modeling in general:
 - all models are wrong; some are useful
 - even if all models are wrong, we should prefer models that are less wrong in the sense of Isaak
 - so how our (less) wrong models useful?
 - while for data: garbage in — garbage out
 - for models: more assumptions in — more conclusions out
 - e.g., causal modeling
 - reasonable assumptions in; interesting conclusions out
 - all models contain some assumptions about the **data-generating process**
 - more assumptions about the nature of that process → more information gained from the data

4 Towards the cognitive turn

References

- Austin, John L. (1962). *How to do Things with Words: The William James Lectures delivered at Havard University in 1955*. Oxford: Clarendon.
- Benz, Anton (2006). “Utility and Relevance of Answers”. In: *Game Theory and Pragmatics*. Ed. by Anton Benz, Gerhard Jäger, and Robert van Rooij. Hampshire: Palgrave, pp. 195–219.
- Benz, Anton and Robert van Rooij (2007). “Optimal Assertions and what they Implicate”. In: *Topoi* 26, pp. 63–78.
- Bunt, Harry and William Black (2000). *Abduction, Belief and Context in Dialogue*. Amsterdam, Philadelphia: John Benjamins.
- Cummins, Chris and Napoleon Katsos, eds. (2019). *The Oxford Handbook of Experimental Semantics and Pragmatics*. Oxford: Oxford University Press.
- Degen, Judith (2023). “The Rational Speech Act Framework”. In: *Annual Review of Linguistics* 9.1, pp. 519–540.
- Frank, Michael C. and Noah D. Goodman (2012). “Predicting Pragmatic Reasoning in Language Games”. In: *Science* 336.6084, p. 998.
- Franke, Michael and Judith Degen (2023). “The softmax function: Properties, motivation, and interpretation”. In.
- Franke, Michael and Gerhard Jäger (2016). “Probabilistic pragmatics, or why Bayes’ rule is probably important for pragmatics”. In: *Zeitschrift für Sprachwissenschaft* 35.1, pp. 3–44.
- Gazdar, Gerald (1979). *Pragmatics: Implicature, Presupposition, and Logical Form*. New York: Academic Press.
- Grice, Paul Herbert (1975). “Logic and Conversation”. In: *Syntax and Semantics, Vol. 3, Speech Acts*. Ed. by Peter Cole and Jerry L. Morgan. New York: Academic Press, pp. 41–58.
- Hobbs, Jerry R., Mark Stickel, and Paul Martin (1993). “Interpretation as Abduction”. In: *Artificial Intelligence* 63, pp. 69–142.
- Kadmon, Nirit (2001). *Formal Pragmatics*. Oxford: Wiley-Blackwell.
- Karttunen, Lauri (1973). “Presuppositions of Complex Sentences”. In: *Linguistic Inquiry* 4, pp. 167–193.
- Meibauer, Jörg and Markus Steinbach, eds. (2011). *Experimental Pragmatics/Semantics*. Amsterdam, Philadelphia: John Benjamins.
- Merin, Arthur (1994). “Algebra of Elementary Social Acts”. In: *Foundations of Speech Act Theory. Philosophical and Linguistic Perspectives*. Ed. by Savas L. Tsohatzidis. London and New York: Routledge, pp. 234–263.
- (1999). “Information, Relevance, and Social Decisionmaking: Some Principles and Results of Decision-Theoretic Semantics”. In: *Logic, Language and Computation*. Ed. by Lawrence C. Moss, Jonathan Ginzburg, and Maarten de Rijke. Vol. 2. CSLI Publications, pp. 179–221.
- Noveck, Ira A. (2018). *Experimental Pragmatics: The Making of a Cognitive Science*. New York: Cambridge University Press.
- Noveck, Ira A. and Dan Sperber, eds. (2004). *Experimental Pragmatics*. Hampshire: Palgrave MacMillan.
- Russell, Benjamin (2012). “Probabilistic Reasoning and the Computation of Scalar Implicatures”. PhD thesis. Brown University.
- Scontras, Gregory, Michael Henry Tessler, and Michael Franke (2021). *A practical introduction to the Rational Speech Act modeling framework*.

- Searle, John R. (1969). *Speech Acts: An Essay in the Philosophy of Language*. Cambridge University Press.
- Sprouse, Jon (2007). “A program for experimental syntax: Finding the relationship between acceptability and grammatical knowledge”. PhD thesis. University of Maryland.
- ed. (2023). *The Oxford Handbook of Experimental Syntax*. Oxford University Press.
- Stalnaker, Robert (1973). “Presupposition”. In: *The Journal of Philosophical Logic* 2, pp. 447–457.
- Zeevat, Henk (2014). *Language Production and Interpretation: Linguistics meets Cognition*. Current Research in the Semantics/Pragmatics Interface. Leiden: Brill.